

Fine-Tuning Transformer Models for NLP Tasks

(Sentiment Analysis)

Team Members

- Ahmed Abdelaziz Hareedy
- Anas Ahmed Desoky
- Salah Amer Mohamed
- Menna Mohamed Abdelhady
- Horia Ahmed Abdelatief
- Mayar Mohamed Khedr

Under Supervision:

Eng. Maryam Mahmoud

1. Problem Statement

The goal is to build a sentiment analysis model that accurately classifies IMDB movie reviews into **positive** or **negative** sentiments. This task is valuable in understanding audience feedback for content creators and businesses.

2. Dataset and Preprocessing

Dataset Used: IMDB 50K Movie Reviews

- 50,000 reviews, equally balanced (25k positive / 25k negative)
- Text + Sentiment (positive or negative)

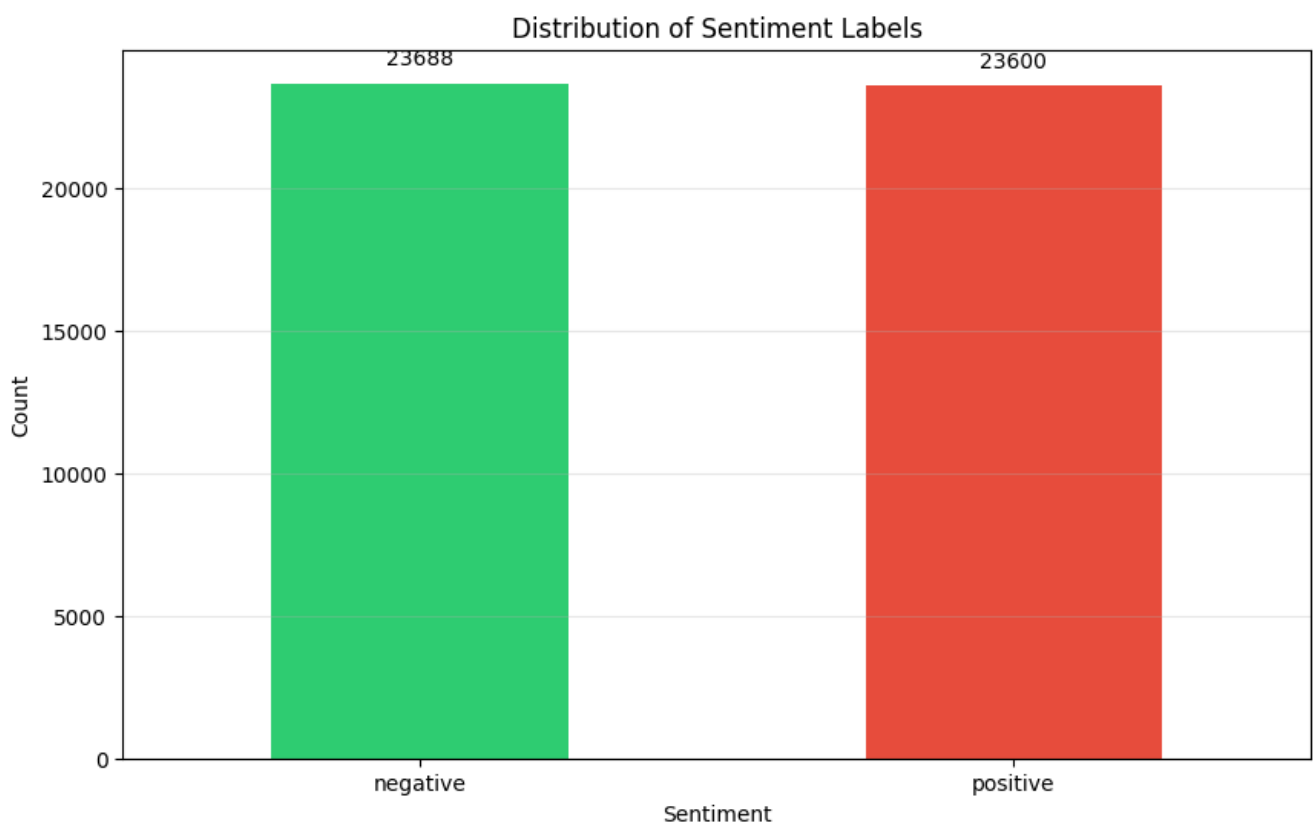
Preprocessing Steps:

- **HTML Tag Removal:** Strips out <div>,

- **Lowercasing:** Converts all characters to lowercase to avoid duplication (Movie, movie → movie).
- **URL, Mention & Hashtag Removal:** Eliminates irrelevant patterns like links and Twitter-style tags.
- **Punctuation & Special Characters:** Removed using regular expressions.
- **Tokenization:** Splits sentences into words.
- **Stopword Removal:**
 - **First:** Removed all Common English stopwords (“the”, “is”, “not”).
 - **Second in Optimization:** Remove all Common stopwords except the negative words like (not, nor, shouldn't)

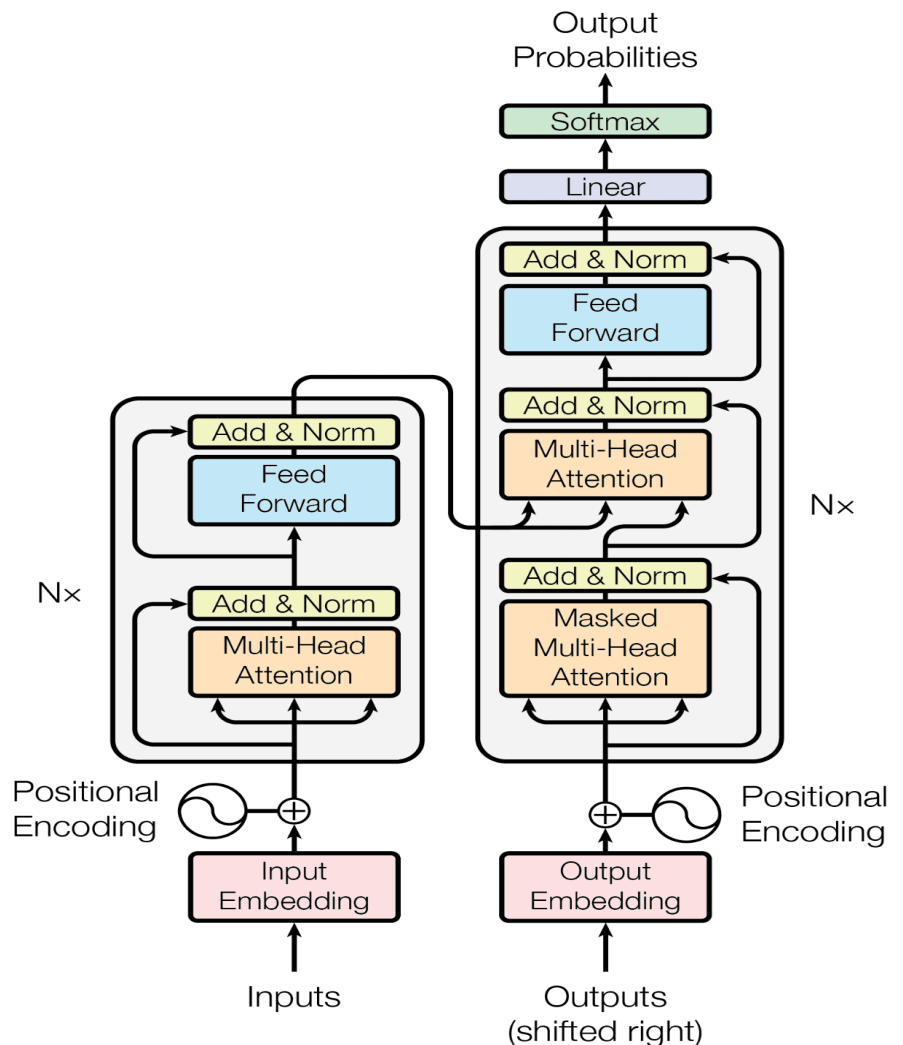
- **Lemmatization** (**Use before Optimization**) Reduces words to their base form (running → run).
- **Encoding Label:** Sentiment labels were converted to numeric form (positive → 1, negative → 0).

Dataset After Preprocessing



➤ Solve our problem using Transformers

Transformers Architecture



3. Models Selection Rationale

For this sentiment analysis task, we used two transformer-based models:

- **DistilBERT:** A distilled version of BERT, chosen for its significantly reduced size and faster inference time while retaining over 95% of BERT's performance. It is ideal for real-time or resource-limited environments.
- **RoBERTa (Robustly Optimized BERT Approach):** Selected for its enhanced training methodology, including more data,

longer sequences, and dynamic masking, resulting in state-of-the-art performance on multiple NLP benchmarks. It is ideal when accuracy is prioritized over speed.

These models were selected to compare the trade-offs between efficiency and performance in transformer-based sentiment classification.

4. Implementation Detail:

- **Pretrained Models:** HuggingFace Transformers library was used to load distilbert-base-uncased and roberta-base.
- **Tokenization:** Used the corresponding tokenizer (DistilBertTokenizerFast and RobertaTokenizerFast) to convert raw text into model-compatible input IDs and attention masks.
- **Fine-Tuning:**
 - Both models were fine-tuned on the preprocessed dataset.
 - Used AdamW optimizer and CrossEntropyLoss.
 - Training was conducted over 5–10 epochs with early stopping based on validation loss.
- **Train/Test Split:** The dataset was split 70% Training, 15% Validation, 15% Test

Optimize RoBERTa Training Details:

Using following:


- Cosine Learning Rate Scheduler with Warmup
- Weight Decay in Optimizer : helps reduce overfitting.
- Early Stopping
- Increase number of Epochs


Train/Test Split: The dataset was split 70% Training, 21% Validation, 9% Test

5. Results and Analysis


Metric	DistilBERT	RoBERTa	Optimized RoBERTa
Accuracy	89.83%	90.51%	93.27%
Precision	89.83%	90.51%	93.28%
Recall	89.83%	90.51%	93.28%
F1-Score	89.83%	90.51%	93.26%
Avg. Inference Time	0.0042 s	~0.005–0.006 s	~0.0055 s


RoBERTa Analysis Before Optimization





RoBERTa Analysis After Optimization





RoBERTa Prediction



```
# Example usage
sample_texts = [
    "I absolutely love this product! It's amazing!",
    "This is the worst experience I've ever had."
]

for text in sample_texts:
    sentiment = predict_sentiment(text, model, tokenizer, device)
    print(f"Text: {text}")
    print(f"Sentiment: {sentiment}\n")
```

Text: I absolutely love this product! It's amazing!
Sentiment: positive

Text: This is the worst experience I've ever had.
Sentiment: Negative

DistilBERT Analysis



DistilBERT Prediction

```
tokenizer = DistilBertTokenizer.from_pretrained('sentiment_model')
for text in sample_texts:
    sentiment = predict_sentiment(text, model, tokenizer)
    print(f"Text: {text}")
    print(f"Sentiment: {sentiment}\n")
```

Python

Sample Predictions:

/tmp/ipykernel_31/3517145775.py:46: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses t
model.load_state_dict(torch.load('best_model.pt'))

Text: I absolutely love this product! It's amazing!

Sentiment: positive

Text: This is the worst experience I've ever had.

Sentiment: negative

Text: The movie was fantastic, with great acting and a compelling story.

Sentiment: positive

Text: I was really disappointed with the service at this restaurant.

Sentiment: negative



Model Api Result

Sentiment Analysis

Enter your text to analyze its sentiment
(positive/negative).

Enter your text

I absolutely love this product! It's amazing!

Analyze Sentiment

Analysis Result:

Sentiment: Positive

Confidence: 99.79%

Sentiment Analysis

Enter your text to analyze its sentiment
(positive/negative).

Enter your text

This is the worst experience I've ever had."

Analyze Sentiment

Analysis Result:

Sentiment: Negative

Confidence: 82.67%